

Constrained Optimization in RL

- Constrained Optimization in RL
- Basic knowledge
 - Descent direction
 - Stepsize
 - Project Subgradient
 - Proximal Algorithms
 - Extra-gradient
 - Natural Gradient
 - Conjugate gradient
 - Bregman Divergence and Mirror Descent
- Constrained MDP
 - Trust-region-based
 - CPO: Constrained Policy Optimization (ICML 2017)
 - PCPO: Projection-Based Constrained Policy Optimization (ICLR 2020)
 - FOCOPS: First Order Constrained Optimization in Policy Space (NIPS 2020)
 - Lyapunov Optimization in Stochastic Networks
 - SDQN: A Lyapunov-based Approach to Safe Reinforcement Learning (NIPS 2018)
 - SPG: Lyapunov-based Safe Policy Optimization for Continuous Control (ICML 2019)
 - LBPO: Lyapunov Barrier Policy Optimization (2021)
 - Primal-Dual Optimization
 - CRPO: A New Approach for Safe Reinforcement Learning with Convergence Guarantee
 - PDO: Risk-Constrained Reinforcement Learning with Percentile Risk Criteria
 - RCPO: Reward Constrained Policy Optimization
 - OPDOP: Provably Efficient Safe Exploration via Primal-Dual Policy Optimization
 - CPPO: Responsive Safety in Reinforcement Learning by PID Lagrangian Methods
 - Convergent Policy Optimization for Safe Reinforcement Learning (NIPS 2019)
 - Safe layer
 - Safe Exploration in Continuous Action Spaces
 - Evolutionary approach
 - Constrained Cross-Entropy Method for Safe Reinforcement Learning (NIPS 2018)

Basic knowledge

Descent direction

1. Lemma 1: Given $x \in \mathbb{R}^n$ and $f : \mathbb{R}^n \rightarrow \mathbb{R}$, among all directions from x , the direction $d = -\nabla f(x)$ gives the maximum rate of decrease in terms of the value of f .
2. Lemma 2: Given $x \in \mathbb{R}^n$ and $f : \mathbb{R}^n \rightarrow \mathbb{R}$, any direction $d \in \mathbb{R}^n$ satisfying $\langle d, \nabla f(x) \rangle < 0$ is a descent direction.
3. Lemma 3: Given $x \in \mathbb{R}^n$, $f : \mathbb{R}^n \rightarrow \mathbb{R}$, $S \in \mathbb{S}_{++}^n$, the direction $-S\nabla f(x) \in \mathbb{R}^n$ is a descent direction at x .

4. Different local approximation:

1. $f(x) \approx f(x_k) + \nabla f(x_k)^\top (x - x_k) + \frac{1}{2\alpha_k} (x - x_k)^\top S^{-1} (x - x_k)$
2. $x_{k+1} = x_k - \alpha_k \cdot S_k \nabla f(x_k)$, Where $S_k \in \mathbb{S}_{++}^n$
 1. Gradient descent: $S_k = I_n$
 2. Newton direction: $S_k = [\nabla^2 f(x_k)]^{-1}$
 3. Scaled Newton direction: $S_k = \text{diag}([\nabla^2 f(x_k)]^{-1})$
 4. Quasi-Newton direction: approximation of $[\nabla^2 f(x_k)]^{-1}$
 5. Regularized Newton direction: $S_k = [\nabla^2 f(x_k) + \mu_k I]^{-1}$

Stepsize

1. Assumption 1: L-Lipschitz continuous gradient: $\|\nabla f(x) - \nabla f(y)\|_2 \leq L\|x - y\|_2$
2. Lemma 4: if f has L-Lipschitz continuous gradient and convex, then (s and n) $f(y) \leq f(x) + \nabla f(x)^\top (y - x) + \frac{L}{2}\|y - x\|_2^2$.
3. $\alpha = \frac{1}{L}$, then $f(x_k) - f^* \leq \frac{\|x_0 - x^*\|_2^2}{2\alpha k}$

Project Subgradient

1. First-order condition for convex function: $f(y) \geq f(x) + \nabla f(x)^\top (y - x)$
2. Subgradient: $f(x) \geq f(x_0) + y^\top (x - x_0)$
 1. $y \in \mathbb{R}^n$ is called a subgradient
 2. The set $\partial f(x_0)$ of all subgradients is the subdifferential of f at x_0
3. Supporting hyperplane
 1. Definition 3 (Supporting hyperplane). Consider a nonempty set $C \subseteq \mathbb{R}^n$ and a boundary point $x_0 \in \text{bd}(C)$. If $a \neq 0$ in \mathbb{R}^n satisfies $a^\top x \leq a^\top x_0, \forall x \in C$, then $\{x \mid a^\top x = a^\top x_0\}$ is called a supporting hyperplane to C at x_0 .
 2. Theorem 1 (Weak Separating Hyperplane Theorem). Consider any convex set $C \subseteq \mathbb{R}^n$ and a point $x_0 \in \mathbb{R}^n / C$. Then, there exist $a \neq 0$ (in \mathbb{R}^n) and $b \in \mathbb{R}$ with $a^\top x \leq b$ and $a^\top x_0 \geq b, \forall x \in C$
3. For convex function, $\partial f(x_0) \neq \emptyset$

3. Orthogonal Projection

1. $P_C(x) \equiv \arg \min_{y \in C} \|y - x\|$
2. The projection $y^* \in C$ is unique
3. non-expansive:
 1. $P_C(x) - P_C(y) \leq \|x - y\|$
4. Projection gradient
 1. $x_{k+1} = P_C(x_k - \alpha_k \nabla f(x_k))$
 2. equal to optimize the local approximation:
 1. $x_{k+1} = \arg \min_{x \in C} \{f(x_k) + \nabla_k f(x_k)^\top (x - x_k) + \frac{L}{2}\|x - x_k\|^2\}$

Proximal Algorithms

1. Proximal operator: $\text{prox}_f(x) = \arg \min_{u \in \mathbb{R}^n} \{f(u) + \frac{1}{2}\|u - x\|^2\}$
 1. Consider optimizing the composite function $f(x) + g(x)$

1. Project gradient:

$$\begin{aligned}
 x_{k+1} &= \arg \min_{x \in \mathbb{R}^n} \left\{ f(x_k) + \nabla f(x_k)^\top (x - x_k) + g(x) + \frac{1}{2\alpha_k} \|x - x_k\|^2 \right\} \\
 &= \arg \min_{x \in \mathbb{R}^n} \left\{ \alpha_k g(x) + \frac{1}{2} \|x - (x_k - \alpha_k \nabla f(x_k))\|^2 \right\} \\
 &= \text{prox}_{\alpha_k g}(x_k - \alpha_k \nabla f(x_k))
 \end{aligned}$$

2. gradient descent for $f(x_k)$, then project to $f(x_k) + g(x)$

2. Proximal gradient descent

$$1. T_L^{f,g}\{x\} = \text{prox}_{\frac{1}{L}g}\left(x - \frac{1}{L}\nabla f(x)\right)$$

3. The Augment Lagrangian Methods

$$1. H^* = \min \{H(x, z) \equiv h_1(x) + h_2(z) \mid Ax + Bz = c\}$$

2. Lagrangian function:

$$\begin{aligned}
 L(x, z, y) &= h_1(x) + h_2(z) + y^\top (Ax + Bz - c) \\
 &= h_1(x) + y^\top Ax + h_2(z) + y^\top Bz - y^\top c
 \end{aligned}$$

3. dual function: $g(y) = \min_{x,z} \{h_1(x) + y^\top Ax + h_2(z) + y^\top Bz - y^\top c\}$

1. minimize the $-g(y)$ using proximal methods:

$$1. y_{k+1} = \arg \min_y \{-g(y) + \frac{1}{2\rho} \|y - y_k\|^2\}$$

2. stationary point: $y_{k+1} = y_k + \rho \nabla_y g(y_{k+1})$

$$1. y_{k+1} = y_k + \rho(A^T x_{k+1} + B^T z_{k+1} - c)$$

$$2. x_{k+1} = \arg \min_x h_1(x) + y_{k+1}^T Ax$$

$$1. \text{stationary point: } 0 = A^T(y_k + \rho(A^T x_{k+1} + B^T z_{k+1} - c)) + \partial_x h_1(x_{k+1})$$

$$3. z_{k+1} = \arg \min_z h_2(z) + y_{k+1}^T Bz$$

$$1. \text{stationary point: } 0 = B^T(y_k + \rho(A^T x_{k+1} + B^T z_{k+1} - c)) + \partial_z h_2(z_{k+1})$$

3. x_{k+1}, z_{k+1} equal to solve:

$$1. H(x, z) = h_1(x) + h_2(z) + \frac{\rho}{2} \|Ax + Bz - c + \frac{1}{\rho} y_k\|^2$$

4. algorithms:

$$1. x_{k+1}, z_{k+1} = \arg \min_{x,z} H(x, z)$$

$$2. y_{k+1} = y_k + \rho(A^T x_{k+1} + B^T z_{k+1} - c)$$

5. ADMM

$$1. x_{k+1} = \arg \min_x H(x, z_k)$$

$$2. z_{k+1} = \arg \min_z H(x_{k+1}, z)$$

$$3. y_{k+1} = y_k + \rho(A^T x_{k+1} + B^T z_{k+1} - c)$$

Extra-gradient

1. projection gradient descent:

$$1. x_{k+1} = P_C(x_k - \alpha_k \nabla_x f(x_k))$$

$$2. \text{optimality condition: } x^* = P_C(x^* - \alpha_k \nabla_x f(x^*))$$

2. Extrapolation

$$1. \bar{x}_k = x_k + \beta(x_k - x_{k-1})$$

2. $x_{k+1} = x_k - \alpha \nabla_x f(\bar{x}_k)$
3. Project extra-gradient:
 1. $\bar{x}_{k+1} = P_C(x_k - \alpha_k \nabla_x f(x_k))$
 2. $x_{k+1} = P_C(x_k - \alpha_k \nabla_x f(\bar{x}_{k+1}))$

Natural Gradient

1. When optimizing distribution: $p(x|\theta)$
 1. distance between different distribution: $KL[p|q] = \int p(x) \log \frac{p(x)}{q(x)}$
 2. local approximation of KL divergence:
 1. $KL[p(x|\theta)|p(x|\theta + \delta)] \approx \int p(x|\theta)[\log p(x|\theta) - (\log p(x|\theta) + \delta \nabla \log p(x|\theta) + \frac{1}{2} \delta^2 \nabla^2 \log p(x|\theta))]$
 2. $KL \approx \frac{1}{2} \delta E_p[\nabla^2 \log p(x|\theta)] \delta$
 3. $H = E_p[\nabla^2 \log p(x|\theta)] = E_p[\nabla \log p(x|\theta) \nabla \log p(x|\theta)^T]$
 3. Descent direction in the trust region:
 1. $\min_{D_{KL}[p_\theta || \theta + \delta\theta] \leq \epsilon} L(\theta + \delta\theta)$
 2. Lagrangian: $\min L(\theta + \delta\theta) + \lambda (D_{KL}[p(\theta) || p_{\theta + \delta\theta}] - \epsilon)$
 3. Approximation: $L(\theta) + \nabla L(\theta)^T \delta\theta + \lambda (\frac{1}{2} \delta\theta^T H \delta\theta - \epsilon)$
 4. Direction: $\delta\theta^* = -\frac{1}{\lambda} H^{-1} \nabla_\theta L(\theta)$

Conjugate gradient

1. Gradient descent
 1. exact step size search
 1. residue: $r_k = b - Qx_k$
 2. stepsize: $\alpha_k = \frac{r_k^T r_k}{r_k^T Q r_k}$
 3. update: $x_{k+1} = x_k + \alpha_k + r_k$
 2. When Q is ill-conditioned, converge slowly.
2. Conjugate
 1. $x_i Q x_j = 0$, then x_i, x_j are conjugate vectors of $Q \in S^+$
 1. $Q = I \rightarrow x_i x_j = 0$, orthogonal is a special case of conjugate
 2. only n independent conjugate vectors for Q
3. Gram-Schmidt Orthogonalization
 1. use n independent basis vector $\{a_i\}$ to construct orthogonal basis $\{q_i\}$
 1. $q_k = a_k + \sum_{j=1}^k b_{kj} q_j$,
 2. construct k using basis a_i with $i \leq k$
 3. create sequential dependence.
 2. weights $b_k = -\sum_{i=1}^{k-1} \frac{\langle a_k, q_i \rangle}{\langle q_i, q_i \rangle}$
 1. proof by inner product: $\langle q_k, q_i \rangle = \langle a_k, q_i \rangle + \sum_{j=1}^k b_{kj} \langle q_j, q_i \rangle$
 2. $\langle q_k, q_i \rangle = \langle a_k, q_j \rangle + b_{ki} \langle q_i, q_i \rangle$
4. Solving KKT condition: $\nabla f(x) = Qx - b = 0$
 1. find n conjugate vectors $p_{i=1}^m$ to combine $x^* = x_0 + \sum_{i=1}^m \alpha_i p_i$
 2. Given initial point x_0 and direction $p_k = \nabla f(x_0) = b - Qx_0$

3. line-search α_k . in linear case, $\alpha_k = \frac{r_k^T r_k}{r_k^T Q r_k}$
4. update descent direction through Gram-Schmidt Orthogonalization
 1. $r_k = \nabla f(x_k) = b - Qx_k$
 - 2.
5. Next point $x_{k+1} = x_k + \alpha_k p_k$
6. $x_{k+1} = x_k + \alpha p_k \rightarrow Qx_{k+1} - b = Qx_k - b + \alpha p_k$
 1. $\nabla f(x_{k+1}) = \nabla f(x_k) + \alpha Q p_k$

Bregman Divergence and Mirror Descent

1. Bregman Divergence

1. Generalize squared Euclidean distance
2. Definition 1 (Bregman divergence) Let $\psi : \Omega \rightarrow \mathbb{R}$ be a function that is: a) strictly convex, b) continuously differentiable, c) defined on a closed convex set Ω . Then the Bregman divergence is defined as

$$\Delta_\psi(x, y) = \psi(x) - \psi(y) - \langle \nabla \psi(y), x - y \rangle, \quad \forall x, y \in \Omega$$

That is, the difference between the value of ψ at x and the first order Taylor expansion of ψ around y evaluated at point x .

3. examples:

1. Euclidean distance: $\psi(x) = \frac{1}{2} \|x\|^2$
2. KL divergence: $\psi(x) = \sum_i x_i \log x_i$
3. L_p norm: $\psi(x) = \frac{1}{2} \|x\|_q^2, \frac{1}{p} + \frac{1}{q} = 1$
4. strong convex case: $\psi(x) \geq \psi(y) + \langle \nabla \psi(y), x - y \rangle + \frac{\sigma}{2} \|x - y\|^2$
 1. $\Delta_\psi(x, y) \geq \frac{\sigma}{2} \|x - y\|^2$

4. Property:

1. Strict convexity
2. Non-negativity
3. Asymmetry
4. Generalized triangle inequality
5. gradient: $\frac{\partial}{\partial x} \Delta_\psi(x, y) = \nabla \psi(x) - \nabla \psi(y)$
5. Projection: $x^* = \operatorname{argmin}_{x \in C} \Delta_\psi(x, x_0)$
 1. Pythagorean Theorem: $\Delta_\psi(y, x_0) \geq \Delta_\psi(y, x^*) + \Delta_\psi(x^*, x_0)$
 2. Proximal operator with Bregman Divergence $x^* = \operatorname{argmin}_{x \in C} \{L(x) + \Delta_\psi(x^*, x_0)\}$
 1. If $L(x)$ is convex:
 1. $L(y) + \Delta_\psi(y, x_0) \geq L(x^*) + \Delta_\psi(x^*, x_0) + \Delta_\psi(y, x^*)$

2. Mirror Descent:

1. Local approximation under L-2 distance:
 1. $f(x) \approx f(x_k) + \nabla f(x_k)^\top (x - x_k) + \frac{1}{2\alpha_k} (x - x_k)^\top S^{-1} (x - x_k)$
 2. Gradient descent: $-\alpha_k \nabla f(x_k)$
2. Approximation with Bregman divergence

$$1. x_{k+1} = \underset{x \in C}{\operatorname{argmin}} \left\{ f(x_k) + \langle g_k, x - x_k \rangle + \frac{1}{\alpha_k} \Delta_\psi(x, x_k) \right\}$$

2. unconstrained case:

$$1. x_{k+1} = (\nabla \psi)^{-1} (\nabla \psi(x_k) - \alpha_k g_k)$$

Constrained MDP

1. Maximize reward: $J^R(\pi) \doteq \mathbb{E}_{\tau \sim \pi} [\sum_{t=0}^{\infty} \gamma^t R(s_t, a_t)]$
2. Constrained by total cost of constraints violation: $J^C(\pi) \doteq \mathbb{E}_{\tau \sim \pi} [\sum_{t=0}^{\infty} \gamma^t C(s_t, a_t)] \leq h$
3. Policy improvement theorem: $J^R(\pi') - J^R(\pi) = \frac{1}{1-\gamma} \mathbb{E}_{\substack{s \sim d^{\pi'} \\ a \sim \pi'}} [A_R^\pi(s, a)]$
4. Solving framework
 1. Linear programming
 2. Lagrangian methods
 1. primal-dual methods
 3. Trust region optimization
 1. CPO
 4. Lyapunov functions

Trust-region-based

1. bound for policy update:
 1. for non-parametric moving average policy:
 1. $\eta(\pi_{\text{new}}) \leq L_{\pi_{\text{old}}}(\pi_{\text{new}}) + \frac{2\epsilon\gamma}{(1-\gamma)^2} \alpha^2$
 2. $\epsilon = \epsilon = \max_s |\mathbb{E}_{a \sim \pi'(a|s)} [A_\pi(s, a)]|$
 3. $\pi_{\text{new}}(a | s) = (1 - \alpha)\pi_{\text{old}}(a | s) + \alpha\pi'(a | s)$
 4. consider the change of state visitation probability
 2. for parametric policy with KL divergence bounded:
 1. $\eta(\tilde{\pi}) \leq L_\pi(\tilde{\pi}) + CD_{\text{KL}}^{\max}(\pi, \tilde{\pi})$, where $C = \frac{2\epsilon\gamma}{(1-\gamma)^2}$
 2. $D_{\text{TV}}^{\max}(\pi, \tilde{\pi}) = \max_s D_{\text{TV}}(\pi(\cdot | s) \| \tilde{\pi}(\cdot | s))$
 3. ignore the change of state visitation probability
2. Approximation

$$\begin{aligned} & \underset{\theta}{\text{maximize}} L_{\theta_{\text{old}}}(\theta) \\ & \text{subject to } \bar{D}_{\text{KL}}^{\rho_{\theta_{\text{old}}}}(\theta_{\text{old}}, \theta) \leq \delta \end{aligned}$$

1. max KL divergence -> mean KL divergence for sampling
2. sample-based average
3. importance sampling to reuse samples

$$\begin{aligned} & \underset{\theta}{\text{minimize}} \mathbb{E}_{s \sim \rho_{\theta_{\text{old}}}, a \sim q} \left[\frac{\pi_\theta(a|s)}{q(a|s)} Q_{\theta_{\text{old}}}(s, a) \right] \\ & \text{subject to } \mathbb{E}_{s \sim \rho_{\theta_{\text{old}}}} [D_{\text{KL}}(\pi_{\theta_{\text{old}}}(\cdot | s) \| \pi_\theta(\cdot | s))] \leq \delta \end{aligned}$$

4. first-order approximation of objective function
 1. $[\nabla_\theta L_{\theta_{\text{old}}}(\theta)|_{\theta=\theta_{\text{old}}} \cdot (\theta - \theta_{\text{old}})]$

5. second-order approximation of constraint

1. KL-divergence

1. Fish-information matrix: H

$$2. \frac{1}{2} \delta \|\theta - \theta_{\text{old}}\|^T H \|\theta - \theta_{\text{old}}\| \leq \delta$$

3. compute the H^{-1} using conjugate gradient

2. L-2 distance

$$1. \frac{1}{2} \|\theta - \theta_{\text{old}}\|^2 \leq \delta$$

CPO: Constrained Policy Optimization (ICML 2017)

1. Joint optimization

$$\begin{aligned} \pi_{k+1} &= \arg \max_{\pi \in \Pi_{\theta}} \mathbb{E}_{s \sim d^{\pi_k}} [A^{\pi_k}(s, a)] \\ \text{s.t. } J_{C_i}(\pi_k) + \frac{1}{1-\gamma} \mathbb{E}_{s \sim d^{\pi_k}} \mathbb{E}_{a \sim \pi} [A_{C_i}^{\pi_k}(s, a)] &\leq d_i \quad \forall i \\ \bar{D}_{KL}(\pi \| \pi_k) &\leq \delta. \end{aligned}$$

1. approximation as TRPO

1. first-order approximation of objective function

2. first-order approximation of cost constraint

3. second-order approximation of KL divergence

2. backtracking line search is used to ensure surrogate constraint satisfaction

PCPO: Projection-Based Constrained Policy Optimization (ICLR 2020)

1. Two-step algorithm

1. performs a local reward improvement update

2. projecting the policy back onto the constraint set

2. Step 1: trust region policy optimization

$$\begin{aligned} \pi^{k+\frac{1}{2}} &= \arg \max_{\pi} \mathbb{E}_{s \sim d^{\pi}} [A_R^{\pi^k}(s, a)] \\ \text{s.t. } \mathbb{E}_{s \sim d^{\pi^k}} [D_{KL}(\pi \| \pi^k)[s]] &\leq \delta \end{aligned}$$

1. approximation as before

3. Step 2: constraint-satisfying projection

$$\begin{aligned} \pi^{k+1} &= \arg \min_{\pi} D(\pi, \pi^{k+\frac{1}{2}}) \\ \text{s.t. } J^C(\pi^k) + \mathbb{E}_{s \sim d^{\pi_k}} [A_C^{\pi^k}(s, a)] &\leq h \end{aligned}$$

1. approximation as before

FOCOPS: First Order Constrained Optimization in Policy Space (NIPS 2020)

$$\begin{aligned}
& \underset{\pi_\theta \in \Pi_\theta}{\text{maximize}} && \mathbb{E}_{\substack{s \sim d^{\pi_{\theta_k}} \\ a \sim \pi_\theta}} [A^{\pi_{\theta_k}}(s, a)] \\
& \text{subject to} && J_C(\pi_{\theta_k}) + \frac{1}{1-\gamma} \mathbb{E}_{\substack{s \sim d^{\pi_{\theta_k}} \\ a \sim \pi_\theta}} [A_C^{\pi_{\theta_k}}(s, a)] \leq b \\
& && \bar{D}_{\text{KL}}(\pi_\theta \parallel \pi_{\theta_k}) \leq \delta.
\end{aligned}$$

1. Errors in CPO

1. Sampling error
2. Approximation error
3. conjugate gradient error

2. Two-step algorithm (use formulation in CPO)

1. Solve in nonparameterized policy space

$$\pi^*(a|s) = \frac{\pi_{\theta_k}(a|s)}{Z_{\lambda, \nu}(s)} \exp \left(\frac{1}{\lambda} \left(A^{\pi_{\theta_k}}(s, a) - \nu A_C^{\pi_{\theta_k}}(s, a) \right) \right)$$

1. Z : normalization constant

2. dual variable λ, ν solved by dual function:

$$1. \min_{\lambda, \nu \geq 0} \lambda \delta + \nu \tilde{b} + \lambda \mathbb{E}_{\substack{s \sim d^{\pi_{\theta_k}} \\ a \sim \pi^*}} [\log Z_{\lambda, \nu}(s)]$$

2. Project back into the parameterized policy space

$$\mathcal{L}(\theta) = \mathbb{E}_{s \sim d^{\pi_{\theta_k}}} [D_{\text{KL}}(\pi_\theta \parallel \pi^*)(s)]$$

3. Practical Implementation

1. approximate dual variable: $\frac{\partial L(\pi^*, \lambda, \nu)}{\partial \nu} = 0$

Lyapunov Optimization in Stochastic Networks

1. Stochastic optimization problem

1. $\mathcal{P}_2 : \min_{\forall t, \alpha(t) \in \mathcal{A}^m} \lim_{T \rightarrow \infty} \frac{1}{T} \sum_{t=0}^{T-1} \mathbb{E}[p(t)]$
2. s.t. $\lim_{T \rightarrow \infty} \frac{1}{T} \sum_{t=0}^{T-1} \mathbb{E}[y_k(t)] \leq 0, k \in \{1, \dots, K\}$

2. Virtual Queues

1. queue length: $Q_k(t) = \max\{Q_k(t) + y_k(t), 0\}$
 1. $\sum_{t=0}^{T-1} y_k(t) \leq Q_k(T) - Q_k(0) = Q_k(T)$
 2. $Q_k(t)^2 \leq (Q_k(t) + y_k(t))^2$
 3. $\frac{1}{2} \sum_{k=1}^K Q_k(t+1)^2 \leq \frac{1}{2} \sum_{k=1}^K Q_k(t)^2 + \frac{1}{2} \sum_{k=1}^K y_k(t)^2 + \sum_{k=1}^K Q_k(t)y_k(t)$

2. Lyapunov function

1. $L(K) = \frac{1}{2} \sum_{t=1}^K Q(t)^2$
2. $\Delta L(K) = L(K+1) - L(K) \leq \frac{1}{2} \sum_{k=1}^K y_k(t)^2 + \sum_{k=1}^K Q_k(t)y_k(t)$
3. Denote B is the upper bound of $\frac{1}{2} \sum_{k=1}^K y_k(t)^2$
 1. $\Delta L(K) \leq B + \sum_{k=1}^K Q_k(t)y_k(t)$

3. Drift-plus-penalty Algorithm

1. $\mathcal{P}_2 : \min_{\forall t, \alpha(t) \in \mathcal{A}^m} \mathbb{E}[\Delta L(K) + Vp(t)]$
 1. approximation of original problem
 2. upper bound: $B + \sum_{k=1}^K Q_k(t)y_k(t) + Vp(t)$

4. Performance Analysis

1. Average goal: $O(\frac{1}{V})$
2. Average queue: $O(V)$

SDQN: A Lyapunov-based Approach to Safe Reinforcement Learning (NIPS 2018)

1. problem:

Problem OPT: Given an initial state x_0 and a threshold d_0 , solve $\min_{\pi \in \Delta} \{C_\pi(x_0) : \mathcal{D}_\pi(x_0) \leq d_0\}$. If there is a non-empty solution, the optimal policy is denoted by π^* .

1. minimize cost only
2. Lyapunov function
 1. analyze the stability of dynamic systems
 1. tracking the energy that a system continually dissipates
 2. represent abstract quantities in a system
 1. steady-state performance of a Markov process
3. Lyapunov for CMDPs
 1. $\mathcal{L}_{\pi_B}(x_0, d_0)$
 1. transient state: $T_{\pi_B, d}[L](x) \leq L(x), \forall x \in \mathcal{X}'$
 1. contraction mapping
 2. terminal state: $L(x) = 0, \forall x \in \mathcal{X} \setminus \mathcal{X}'$
 1. terminal state with 0 cost
 3. initial state: $L(x_0) \leq d_0$
 1. satisfy the constrain threshold
 2. Relation between cost value function and Lyapunov function
 1. exist ϵ that $L_\epsilon(x) = \mathbb{E} \left[\sum_{t=0}^{T^*-1} d(x_t) + \epsilon(x_t) \mid \pi_B, x \right]$
 2. upper bound of optimal cost value function
3. Solve ϵ
 1. safety condition: $d_0 \geq L_{\tilde{\epsilon}}(x) \geq T_{\pi_B, d}[L_{\tilde{\epsilon}}](x)$
 2. solve linear programming

$$\tilde{\epsilon} \in \arg \max_{\epsilon: \mathcal{X}' \rightarrow \mathbb{R}_{\geq 0}} \left\{ \sum_{x \in \mathcal{X}'} \epsilon(x) : d_0 - \mathcal{D}_{\pi_B}(x_0) \geq \mathbf{1}(x_0)^\top \left(I - \{P(x' \mid x, \pi_B)\}_{x, x' \in \mathcal{X}'} \right)^{-1} \epsilon \right\}$$

4. Safe update:

1. state-action Lyapunov function:
 1. $Q_L(x, a) = d(x) + \tilde{\epsilon}(x) + \sum_{x'} P(x' \mid x, a) L_{\tilde{\epsilon}}(x')$
2. L_{π_B} induced policy set:
 1. $(\pi(\cdot \mid x) - \pi_B(\cdot \mid x))^\top Q_L(x, \cdot) \leq \tilde{\epsilon}(x)$
3. update policy:

1. $\pi'(\cdot | x) \in \arg \min_{\pi \in \Delta} \{ \pi(\cdot | x)^\top Q(x, \cdot) \}$
2. Linear programming

SPG: Lyapunov-based Safe Policy Optimization for Continuous Control (ICML 2019)

1. Safe policy optimization:

1. $\pi'(\cdot | x) \in \arg \min_{\pi \in \Delta} \{ \pi(\cdot | x)^\top Q(x, \cdot) \}$

2. $(\pi(\cdot | x) - \pi_B(\cdot | x))^\top Q_L(x, \cdot) \leq \tilde{\epsilon}(x)$

3. two efficient algorithm

1. θ —projection
2. α —projection

2. θ —projection

1. trust region optimization

$$\begin{aligned} \mathcal{C}'_{\pi_\theta}(x_0; \pi_{\theta_B}) = & \mathcal{C}_{\pi_{\theta_B}}(x_0) + \beta \bar{D}_{\text{KL}}(\theta, \theta_B) + \\ & \mathbb{E}_{x \sim \mu_{\theta_B, x_0}, a \sim \pi_\theta} \left[Q_{V_{\theta_B}}(x, a) - V_{\theta_B}(x) \right] \end{aligned}$$

1. first-order approximation
2. average constraint surrogate

3. α —projection

1. safety layer

1. embed the set of Lyapunov constraints into the policy network
 1. project action under Lyapunov constraints
 2. first-order approximation
 3. KKT condition -> OPT-Net
2. an unconstrained optimization problem

LBPO: Lyapunov Barrier Policy Optimization (2021)

1. problem

$$\max_{\pi \in \mathcal{P}} [J_\pi(s_0)] \text{ s.t } D_\pi(s_0) \leq d_0$$

2. Update policies inside the L_{π_B} induced policy set

1. Q-value Evaluation

2. Safe Policy Improvement

1. $\pi_+(\cdot | s) = \max_{\pi \in \mathcal{P}} J_\pi(s_0)$

2. $\int_{a \in \mathcal{A}} (\pi(a | s) - \pi_B(a | s)) Q_{L_{\pi_B}, \hat{\epsilon}}(s, a) da \leq \hat{\epsilon}(s)$

3. convert constrain as log-barrier function

Primal-Dual Optimization

Primal methods

CRPO: A New Approach for Safe Reinforcement Learning with Convergence Guarantee

1. Primal-approach (The alternating mirror descent SA algorithm)

1. convergence guaranteed
2. Two-step:
 1. policy evaluation:
 1. reward value function
 2. cost value function
 2. policy update and constrain update
 1. if constraint is satisfied: update policy using reward value function
 2. if not satisfied: update policy using cost value function

Primal-dual methods

PDO: Risk-Constrained Reinforcement Learning with Percentile Risk Criteria

1. Lagrangian Approach and Reformulation
 1. primal-dual descent-ascent algorithm
 2. sample average estimation

RCPO: Reward Constrained Policy Optimization

1. handle discounted sum and mean constraints
2. Lagrangian:
 1. $\min_{\lambda \geq 0} \max_{\theta} L(\lambda, \theta) = \min_{\lambda \geq 0} \max_{\theta} [J_R^{\pi_{\theta}} - \lambda \cdot (J_C^{\pi_{\theta}} - \alpha)]$
3. Penalized reward functions
 1. $\hat{r}(\lambda, s, a) \triangleq r(s, a) - \lambda c(s, a)$
 2. update actor and critic using penalized value function
 3. update λ

OPDOP: Provably Efficient Safe Exploration via Primal-Dual Policy Optimization

1. Lagrangian:
 1. $\min_{\lambda \geq 0} \max_{\theta} L(\lambda, \theta) = \min_{\lambda \geq 0} \max_{\theta} [J_R^{\pi_{\theta}} - \lambda \cdot (J_C^{\pi_{\theta}} - \alpha)]$
2. utility function over K episodes

$$\text{Regret}(K) = \sum_{k=1}^K \left(V_{r,1}^{\pi^*}(x_1) - V_{r,1}^{\pi^k}(x_1) \right)$$

$$\text{Violation}(K) = \sum_{k=1}^K \left(b - V_{g,1}^{\pi^k}(x_1) \right)$$

3. Learning process
 1. policy evaluation: Least-Squares Temporal Difference
 2. primal update: KL divergence penalized update
 3. dual update: upper bounded gradient

CPPO: Responsive Safety in Reinforcement Learning by PID Lagrangian Methods

1. Lagrangian approaches are in oscillations and over-shoot

1. apply PID to adjust dual variable

$$\begin{aligned}
6: \quad & \Delta \leftarrow J_C - d \\
7: \quad & \partial \leftarrow (J_C - J_{C,prev})_+ \\
8: \quad & I \leftarrow (I + \Delta)_+ \\
9: \quad & \lambda \leftarrow (K_P \Delta + K_I I + K_D \partial)_+ \\
10: \quad & J_{C,prev} \leftarrow J_C
\end{aligned}$$

Convergent Policy Optimization for Safe Reinforcement Learning (NIPS 2019)

1. Lagrangian:
2. Successive convex relaxation:
 1. Both value and constraint

Safe layer

Safe Exploration in Continuous Action Spaces

1. only for immediate-constraint functions
2. linearization cost function:
 1. $\bar{c}_i(s') \triangleq c_i(s, a) \approx \bar{c}_i(s) + g(s; w_i)^\top a$
3. project action: convex optimization

$$\begin{aligned}
a^* = \arg \min_a & \frac{1}{2} \|a - \mu_\theta(s)\|^2 \\
\text{s.t. } & \bar{c}_i(s) + g(s; w_i)^\top a \leq C_i \forall i \in [K]
\end{aligned}$$

Evolutionary approach

Constrained Cross-Entropy Method for Safe Reinforcement Learning (NIPS 2018)

1. Sampling and sorting